

Table of Contents

- 1 Phase3 Project
 - 1.1 Business Objective
- 2 Dataset
- 3 Import libraries
- 4 EDA
 - 4.1 Data Understanding
 - 4.2 Checking for Null values
 - 4.3 Checking the datatypes
 - 4.3.1 Categorical data
 - 4.3.2 Numeric data
 - 4.3.2.1 Scaling the data
 - 4.4 Checking for duplicate data
- 5 Mapping waterpump distribution
- 6 Encoding
 - 6.1 Label encoding approach
 - 6.2 One-hot encoding
- 7 Building baseline models
 - 7.1 Logistic Regression
 - 7.2 Decision Tree
 - 7.3 KNN model
 - 7.4 Selecting a model
- 8 Decision Tree
 - 8.1 Hyperparameter Tuning
 - 8.1.1 max_depth
 - 8.1.2 min_samples_split
 - 8.1.3 min_samples_leaf
 - 8.1.4 model with optimized parameters
 - 8.2 feature_importance
 - 8.2.1 extracting top_10 features
 - 8.3 Decision Tree with the top10 features
 - 8.3.1 GridSearch CV
- 9 Random Forest with top10 features
- 10 Visualize scores
 - 10.1 Confusion Matrix
- 11 Random Forest with all the features
 - 11.1 Confusion matrix
- 12 Examining the target feature
- 13 Training a model with a balanced data set
 - 13.1 Confusion Matrix

- [14 Random forest model on the balanced dataset](#)
- [15 Next Steps](#)

Phase3 Project

Business Objective

To help the Government of Tanzania monitor the condition of installed waterpumps across the country. Given a set of parameters, the model should be able to predict the status of a waterpump. Status can be as classified as:

1. Functional
2. Functional needs repair
3. non functional

Dataset

Dataset sourced from: <https://www.drivendata.org/competitions/7/pump-it-up-data-mining-the-water-table/page/23/>

Import libraries

```
In [1]: import pandas as pd
import numpy as np
import geopandas as gpd
from shapely.geometry import Point, Polygon
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.preprocessing import LabelEncoder
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report
from sklearn.metrics import f1_score, accuracy_score, precision_score, recall_score
from sklearn.tree import DecisionTreeClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.model_selection import GridSearchCV
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import plot_confusion_matrix
from sklearn.preprocessing import MinMaxScaler
from sklearn.preprocessing import StandardScaler
import warnings
warnings.filterwarnings('ignore')
```


EDA

```
In [2]: # importing dataset
df = pd.read_csv('waterwell.csv')
df.head()
```

Out[2]:

	id	amount_tsh	date_recorded	funder	gps_height	installer	longitude	latitude	wpt_name
0	69572	6000.0	3/14/2011	Roman	1390	Roman	34.938093	-9.856322	none
1	8776	0.0	3/6/2013	Grumeti	1399	GRUMETI	34.698766	-2.147466	Zahanati
2	34310	25.0	2/25/2013	Lottery Club	686	World vision	37.460664	-3.821329	Kwa Mahundi
3	67743	0.0	1/28/2013	Unicef	263	UNICEF	38.486161	-11.155298	Zahanati Ya Nanyumbu
4	19728	0.0	7/13/2011	Action In A	0	Artisan	31.130847	-1.825359	Shuleni

5 rows × 41 columns



In [3]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 59400 entries, 0 to 59399
Data columns (total 41 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   id                                     59400 non-null  int64
1   amount_tsh                           59400 non-null  float64
2   date_recorded                         59400 non-null  object
3   funder                                55765 non-null  object
4   gps_height                           59400 non-null  int64
5   installer                             55745 non-null  object
6   longitude                             59400 non-null  float64
7   latitude                             59400 non-null  float64
8   wpt_name                             59400 non-null  object
9   num_private                           59400 non-null  int64
10  basin                                 59400 non-null  object
11  subvillage                           59029 non-null  object
12  region                               59400 non-null  object
13  region_code                           59400 non-null  int64
14  district_code                         59400 non-null  int64
15  lga                                   59400 non-null  object
16  ward                                  59400 non-null  object
17  population                            59400 non-null  int64
18  public_meeting                        56066 non-null  object
19  recorded_by                           59400 non-null  object
20  scheme_management                     55523 non-null  object
21  scheme_name                           31234 non-null  object
22  permit                                56344 non-null  object
23  construction_year                     59400 non-null  int64
24  extraction_type                       59400 non-null  object
25  extraction_type_group                  59400 non-null  object
26  extraction_type_class                  59400 non-null  object
27  management                            59400 non-null  object
28  management_group                      59400 non-null  object
29  payment                               59400 non-null  object
30  payment_type                          59400 non-null  object
31  water_quality                         59400 non-null  object
32  quality_group                         59400 non-null  object
33  quantity                              59400 non-null  object
```

```

34 quantity_group      59400 non-null object
35 source              59400 non-null object
36 source_type         59400 non-null object
37 source_class        59400 non-null object
38 waterpoint_type     59400 non-null object
39 waterpoint_type_group 59400 non-null object
40 status_group        59400 non-null object
dtypes: float64(3), int64(7), object(31)
memory usage: 18.6+ MB

```

Data Understanding

Just a glance at the features, we can see that some features like

latitude, longitude, date_recorded, water_quality, construction_year and so on could be important for modelling. On the other hand, features like wpt_name, scheme_name, lga, ward can be considered as superfluous and can be omitted during modelling. We will make informed decisions based on EDA and feature_importances_ after we build models.

Checking for Null values

```

In [4]: # checking for null values and returning it as a pandas series
empty=df.isna().sum()
empty

```

```

Out[4]: id                0
amount_tsh              0
date_recorded           0
funder                 3635
gps_height              0
installer              3655
longitude               0
latitude                0
wpt_name                0
num_private             0
basin                  0
subvillage             371
region                 0
region_code             0
district_code           0
lga                    0
ward                   0
population              0
public_meeting         3334
recorded_by            0
scheme_management      3877
scheme_name            28166
permit                 3056
construction_year       0
extraction_type         0
extraction_type_group   0
extraction_type_class   0
management              0
management_group        0
payment                 0
payment_type            0
water_quality           0
quality_group           0
quantity                0
quantity_group          0
source                  0

```

```

source_type          0
source_class         0
waterpoint_type      0
waterpoint_type_group 0
status_group         0
dtype: int64

```

```

In [5]: #converting the empty series into a dictionary
empty_dict = dict(empty)

#looping thru dictionary to isolate the columns that have null values
empty_list = []
for key,value in empty_dict.items():
    if value != 0:
        empty_list.append(key)

empty_list
# we now have the list of columns that have null values

```

```

Out[5]: ['funder',
         'installer',
         'subvillage',
         'public_meeting',
         'scheme_management',
         'scheme_name',
         'permit']

```

```

In [6]: # examining those columns
df_empty = df[empty_list]
df_empty

```

```

Out[6]:

```

	funder	installer	subvillage	public_meeting	scheme_management	scheme_name	permit
0	Roman	Roman	Mnyusi B	True	VWC	Roman	False
1	Grumeti	GRUMETI	Nyamara	NaN	Other	NaN	True
2	Lottery Club	World vision	Majengo	True	VWC	Nyumba ya mungu pipe scheme	True
3	Unicef	UNICEF	Mahakamani	True	VWC	NaN	True
4	Action In A	Artisan	Kyanyamisa	True	NaN	NaN	True
...
59395	Germany Republi	CES	Kiduruni	True	Water Board	Losaa Kia water supply	True
59396	Cefa-njombe	Cefa	Igumbilo	True	VWC	Ikondo electrical water sch	True
59397	NaN	NaN	Madungulu	True	VWC	NaN	False
59398	Malec	Musa	Mwinyi	True	VWC	NaN	True
59399	World Bank	World	Kikatanyemba	True	VWC	NaN	True

59400 rows × 7 columns

```
In [7]: df_empty.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 59400 entries, 0 to 59399
Data columns (total 7 columns):
#   Column                Non-Null Count  Dtype
---  -
0   funder                 55765 non-null  object
1   installer              55745 non-null  object
2   subvillage             59029 non-null  object
3   public_meeting         56066 non-null  object
4   scheme_management      55523 non-null  object
5   scheme_name            31234 non-null  object
6   permit                 56344 non-null  object
dtypes: object(7)
memory usage: 3.2+ MB
```

We can see that all the columns that have null values are categorical.
Also, recall from the original df that total number of rows is 59400

```
In [8]: # Looking at scheme_name first since it has the highest number of null values
df['scheme_name'].value_counts()
```

```
Out[8]: K                682
None                644
Borehole            546
Chalinze wate       405
M                   400
...
BL Nkenku kijijini    1
TM Lawate water supply  1
Nguluwater Supply     1
Ujindali              1
BL Laktore            1
Name: scheme_name, Length: 2696, dtype: int64
```

Since 'scheme_name' has approx. 47% of the data missing, even classifying this as 'missing' might skew the analysis. Hence it's best to remove it from the analysis.


```
In [9]: #creating a copy of df for the analysis
df1 = df.copy()
df1.head()
```

```
Out[9]:
```

	id	amount_tsh	date_recorded	funder	gps_height	installer	longitude	latitude	wpt_name
0	69572	6000.0	3/14/2011	Roman	1390	Roman	34.938093	-9.856322	none
1	8776	0.0	3/6/2013	Grumeti	1399	GRUMETI	34.698766	-2.147466	Zahanati
2	34310	25.0	2/25/2013	Lottery Club	686	World vision	37.460664	-3.821329	Kwa Mahundi
3	67743	0.0	1/28/2013	Unicef	263	UNICEF	38.486161	-11.155298	Zahanati Ya Nanyumbu

	id	amount_tsh	date_recorded	funder	gps_height	installer	longitude	latitude	wpt_name
4	19728	0.0	7/13/2011	Action In A	0	Artisan	31.130847	-1.825359	Shuleni

5 rows × 41 columns



```
In [10]: # removing 'scheme_name' from the df
df1.drop('scheme_name',axis=1,inplace=True)
```

Since, the rest of the columns have approx. only 6% of the data missing, we can either choose to drop it or classify it as 'MISSING' for the analysis. Let's classify it as 'MISSING'.

```
In [11]: #replacing the null values as 'MISSING'
df1.fillna('MISSING',inplace=True)
```

```
In [12]: # checking the df
df1.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 59400 entries, 0 to 59399
Data columns (total 40 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   id                                     59400 non-null  int64
1   amount_tsh                           59400 non-null  float64
2   date_recorded                         59400 non-null  object
3   funder                                59400 non-null  object
4   gps_height                            59400 non-null  int64
5   installer                             59400 non-null  object
6   longitude                             59400 non-null  float64
7   latitude                              59400 non-null  float64
8   wpt_name                              59400 non-null  object
9   num_private                           59400 non-null  int64
10  basin                                 59400 non-null  object
11  subvillage                            59400 non-null  object
12  region                                59400 non-null  object
13  region_code                           59400 non-null  int64
14  district_code                         59400 non-null  int64
15  lga                                    59400 non-null  object
16  ward                                   59400 non-null  object
17  population                             59400 non-null  int64
18  public_meeting                        59400 non-null  object
19  recorded_by                           59400 non-null  object
20  scheme_management                     59400 non-null  object
21  permit                                59400 non-null  object
22  construction_year                     59400 non-null  int64
23  extraction_type                       59400 non-null  object
24  extraction_type_group                  59400 non-null  object
25  extraction_type_class                   59400 non-null  object
26  management                             59400 non-null  object
27  management_group                       59400 non-null  object
28  payment                                59400 non-null  object
29  payment_type                           59400 non-null  object
30  water_quality                          59400 non-null  object
31  quality_group                          59400 non-null  object
32  quantity                               59400 non-null  object
33  quantity_group                         59400 non-null  object
```

```

34 source                59400 non-null object
35 source_type           59400 non-null object
36 source_class          59400 non-null object
37 waterpoint_type       59400 non-null object
38 waterpoint_type_group 59400 non-null object
39 status_group          59400 non-null object
dtypes: float64(3), int64(7), object(30)
memory usage: 18.1+ MB

```

We can see that there no more missing values

Checking the datatypes

```
In [13]: # examining the data types of the df
df1.dtypes.value_counts()
```

```
Out[13]: object      30
int64         7
float64        3
dtype: int64
```

We can see that most of the features are categorical

Categorical data

Let's explore some of the features and see if we can glean some information:

```
In [14]: df1[['payment', 'payment_type']]
```

```
Out[14]:
```

	payment	payment_type
0	pay annually	annually
1	never pay	never pay
2	pay per bucket	per bucket
3	never pay	never pay
4	never pay	never pay
...
59395	pay per bucket	per bucket
59396	pay annually	annually
59397	pay monthly	monthly
59398	never pay	never pay
59399	pay when scheme fails	on failure

59400 rows × 2 columns

Since they are the same, we can delete either 'payment' or 'payment_type'

```
In [15]: #dropping 'payment_type' from the df
df1.drop('payment_type',axis=1,inplace=True)
```

```
In [16]: df1[['extraction_type', 'extraction_type_group', 'extraction_type_class']]
```



```
Out[16]:
```

	extraction_type	extraction_type_group	extraction_type_class
0	gravity	gravity	gravity
1	gravity	gravity	gravity
2	gravity	gravity	gravity
3	submersible	submersible	submersible
4	gravity	gravity	gravity
...
59395	gravity	gravity	gravity
59396	gravity	gravity	gravity
59397	swn 80	swn 80	handpump
59398	nira/tanira	nira/tanira	handpump
59399	nira/tanira	nira/tanira	handpump

59400 rows × 3 columns

Again, these are similar and we can chose to eliminate 2 of them from our analysis

```
In [17]: df1.drop(['extraction_type_group', 'extraction_type_class'], axis=1, inplace=True)
```

```
In [18]: df1[['management', 'management_group']]
```

```
Out[18]:
```

	management	management_group
0	vwc	user-group
1	wug	user-group
2	vwc	user-group
3	vwc	user-group
4	other	other
...
59395	water board	user-group
59396	vwc	user-group
59397	vwc	user-group
59398	vwc	user-group
59399	vwc	user-group

59400 rows × 2 columns

```
In [19]: df1['management'].unique()
```

```
Out[19]: array(['vwc', 'wug', 'other', 'private operator', 'water board', 'wua',
```

```
'company', 'water authority', 'parastatal', 'unknown',  
'other - school', 'trust'], dtype=object)
```

```
In [20]: df1['management_group'].unique()
```

```
Out[20]: array(['user-group', 'other', 'commercial', 'parastatal', 'unknown'],  
             dtype=object)
```

We can remove 'management_group' since 'management' provides more detail.

```
In [21]: #drop  
df1.drop('management_group',axis=1,inplace=True)
```

```
In [22]: df1[['source','source_type','source_class']]
```

```
Out[22]:
```

	source	source_type	source_class
0	spring	spring	groundwater
1	rainwater harvesting	rainwater harvesting	surface
2	dam	dam	surface
3	machine dbh	borehole	groundwater
4	rainwater harvesting	rainwater harvesting	surface
...
59395	spring	spring	groundwater
59396	river	river/lake	surface
59397	machine dbh	borehole	groundwater
59398	shallow well	shallow well	groundwater
59399	shallow well	shallow well	groundwater

59400 rows × 3 columns

```
In [23]: df1['source'].unique()
```

```
Out[23]: array(['spring', 'rainwater harvesting', 'dam', 'machine dbh', 'other',  
               'shallow well', 'river', 'hand dtw', 'lake', 'unknown'],  
             dtype=object)
```

```
In [24]: df1['source_type'].unique()
```

```
Out[24]: array(['spring', 'rainwater harvesting', 'dam', 'borehole', 'other',  
               'shallow well', 'river/lake'], dtype=object)
```

```
In [25]: df1['source_class'].unique()
```

```
Out[25]: array(['groundwater', 'surface', 'unknown'], dtype=object)
```

We can choose 'source' over the other two features

```
In [26]: #drop  
df1.drop(['source_type','source_class'],axis=1,inplace=True)
```

```
In [27]: df1[['water_quality','quality_group','quantity','quantity_group','waterpoint_type','wat
```

```
Out[27]:
```

	water_quality	quality_group	quantity	quantity_group	waterpoint_type	waterpoint_type_group
0	soft	good	enough	enough	communal standpipe	communal standpipe
1	soft	good	insufficient	insufficient	communal standpipe	communal standpipe
2	soft	good	enough	enough	communal standpipe multiple	communal standpipe
3	soft	good	dry	dry	communal standpipe multiple	communal standpipe
4	soft	good	seasonal	seasonal	communal standpipe	communal standpipe
...
59395	soft	good	enough	enough	communal standpipe	communal standpipe
59396	soft	good	enough	enough	communal standpipe	communal standpipe
59397	fluoride	fluoride	enough	enough	hand pump	hand pump
59398	soft	good	insufficient	insufficient	hand pump	hand pump
59399	salty	salty	enough	enough	hand pump	hand pump

59400 rows × 6 columns



```
In [28]: df1['water_quality'].unique()
```

```
Out[28]: array(['soft', 'salty', 'milky', 'unknown', 'fluoride', 'coloured',  
              'salty abandoned', 'fluoride abandoned'], dtype=object)
```

```
In [29]: df1['quality_group'].unique()
```

```
Out[29]: array(['good', 'salty', 'milky', 'unknown', 'fluoride', 'colored'],  
              dtype=object)
```

```
In [30]: df1['quantity'].unique()
```

```
Out[30]: array(['enough', 'insufficient', 'dry', 'seasonal', 'unknown'],  
              dtype=object)
```

```
In [31]: df1['quantity_group'].unique()
```

```
Out[31]: array(['enough', 'insufficient', 'dry', 'seasonal', 'unknown'],  
              dtype=object)
```

```
In [32]: df1['waterpoint_type'].unique()
```

```
Out[32]: array(['communal standpipe', 'communal standpipe multiple', 'hand pump',  
              'other', 'improved spring', 'cattle trough', 'dam'], dtype=object)
```

```
In [33]: df1['waterpoint_type_group'].unique()
```

```
Out[33]: array(['communal standpipe', 'hand pump', 'other', 'improved spring',  
              'cattle trough', 'dam'], dtype=object)
```

```
In [34]: #dropping quality_group, quantity_group and waterpoint_type_group  
df1.drop(['quality_group', 'waterpoint_type_group', 'quantity_group'], axis=1, inplace=True)  
df1.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 59400 entries, 0 to 59399  
Data columns (total 31 columns):  
#   Column                Non-Null Count  Dtype  
---  -  
0   id                    59400 non-null  int64  
1   amount_tsh           59400 non-null  float64  
2   date_recorded         59400 non-null  object  
3   funder                59400 non-null  object  
4   gps_height            59400 non-null  int64  
5   installer             59400 non-null  object  
6   longitude             59400 non-null  float64  
7   latitude              59400 non-null  float64  
8   wpt_name              59400 non-null  object  
9   num_private           59400 non-null  int64  
10  basin                 59400 non-null  object  
11  subvillage            59400 non-null  object  
12  region                59400 non-null  object  
13  region_code           59400 non-null  int64  
14  district_code         59400 non-null  int64  
15  lga                   59400 non-null  object  
16  ward                  59400 non-null  object  
17  population            59400 non-null  int64  
18  public_meeting        59400 non-null  object  
19  recorded_by           59400 non-null  object  
20  scheme_management     59400 non-null  object  
21  permit                59400 non-null  object  
22  construction_year     59400 non-null  int64  
23  extraction_type       59400 non-null  object  
24  management            59400 non-null  object  
25  payment               59400 non-null  object  
26  water_quality         59400 non-null  object  
27  quantity              59400 non-null  object  
28  source                59400 non-null  object  
29  waterpoint_type       59400 non-null  object  
30  status_group          59400 non-null  object  
dtypes: float64(3), int64(7), object(21)  
memory usage: 14.0+ MB
```

```
In [35]: df1[['region', 'region_code', 'district_code']]
```

```
Out[35]:
```

	region	region_code	district_code
0	Iringa	11	5
1	Mara	20	2
2	Manyara	21	4
3	Mtwara	90	63
4	Kagera	18	1
...

	region	region_code	district_code
59395	Kilimanjaro	3	5
59396	Iringa	11	4
59397	Mbeya	12	7
59398	Dodoma	1	4
59399	Morogoro	5	2

59400 rows × 3 columns

```
In [36]: df1['public_meeting'].unique()
```

```
Out[36]: array([True, 'MISSING', False], dtype=object)
```

```
In [37]: df1['recorded_by'].unique()
```

```
Out[37]: array(['GeoData Consultants Ltd'], dtype=object)
```

```
In [38]: df1['num_private'].unique()
```

```
Out[38]: array([  0,  39,   5,  45,   6,   3, 698,  32,  15,   7,  25,
        102,   1,  93,  14,  34, 120,  17, 213,  47,   8,  41,
         80, 141,  20,  35, 131,   4,  22,  11,  87,  61,  65,
        136,   2, 180,  38,  62,   9,  16,  23,  42,  24,  12,
        668, 672,  58, 150, 280, 160,  50, 1776,  30,  27,  10,
         94,  26, 450, 240, 755,  60, 111,  300,  55, 1402],
        dtype=int64)
```

```
In [39]: df1[['scheme_management', 'permit']]
```

```
Out[39]:
```

	scheme_management	permit
0	VWC	False
1	Other	True
2	VWC	True
3	VWC	True
4	MISSING	True
...
59395	Water Board	True
59396	VWC	True
59397	VWC	False
59398	VWC	True
59399	VWC	True

59400 rows × 2 columns

```
In [40]: df1['scheme_management'].unique()
```

```
Out[40]: array(['VWC', 'Other', 'MISSING', 'Private operator', 'WUG',
        'Water Board', 'WUA', 'Water authority', 'Company', 'Parastatal',
        'Trust', 'SWC', 'None'], dtype=object)
```

Looks like 'scheme_management' has the same info as 'management' and hence can be removed

```
In [41]: #dropping scheme_management
df1.drop('scheme_management',axis=1,inplace=True)
```

```
In [42]: #dropping id,wpt_name since they will not have a bearing on the analysis
df1.drop(['id','wpt_name'],axis=1,inplace=True)
```

```
In [43]: df1[['funder','installer']]
```

```
Out[43]:
```

	funder	installer
0	Roman	Roman
1	Grumeti	GRUMETI
2	Lottery Club	World vision
3	Unicef	UNICEF
4	Action In A	Artisan
...
59395	Germany Republi	CES
59396	Cefa-njombe	Cefa
59397	MISSING	MISSING
59398	Malec	Musa
59399	World Bank	World

59400 rows × 2 columns

Numeric data

Scaling the data

```
In [44]: # Looking at the distributions of the numerical data
#creating a subset of the numeric data
df1_numeric = pd.DataFrame(df1.select_dtypes(include=['int64','float64']))
df1_numeric.head()
```

```
Out[44]:
```

	amount_tsh	gps_height	longitude	latitude	num_private	region_code	district_code	population
0	6000.0	1390	34.938093	-9.856322	0	11	5	109
1	0.0	1399	34.698766	-2.147466	0	20	2	280
2	25.0	686	37.460664	-3.821329	0	21	4	250
3	0.0	263	38.486161	-11.155298	0	90	63	58

	amount_tsh	gps_height	longitude	latitude	num_private	region_code	district_code	population
4	0.0	0	31.130847	-1.825359	0	18	1	0

```
In [45]: #scaling the numeric data
scaled = MinMaxScaler()
df1_numeric_scaled = pd.DataFrame(data=scaled.fit_transform(df1_numeric),index=df1_numeric.index,
                                   columns=df1_numeric.columns)
df1_numeric_scaled.head()
```

	amount_tsh	gps_height	longitude	latitude	num_private	region_code	district_code	population	c
0	0.017143	0.517483	0.865979	0.153923	0.0	0.102041	0.0625	0.003574	
1	0.000000	0.520629	0.860047	0.815659	0.0	0.193878	0.0250	0.009180	
2	0.000071	0.271329	0.928504	0.671973	0.0	0.204082	0.0500	0.008197	
3	0.000000	0.123427	0.953922	0.042418	0.0	0.908163	0.7875	0.001902	
4	0.000000	0.031469	0.771612	0.843309	0.0	0.173469	0.0125	0.000000	

```
In [46]: # drop the existing numeric data columns
df1.drop(df1_numeric_scaled.columns,axis=1,inplace=True)
```

```
In [47]: #concatenate the numeric df with the scaled values
df1 = pd.concat([df1,df1_numeric_scaled],axis=1,)
df1.head()
```

	date_recorded	funder	installer	basin	subvillage	region	lga	ward	public_me
0	3/14/2011	Roman	Roman	Lake Nyasa	Mnyusi B	Iringa	Ludewa	Mundindi	
1	3/6/2013	Grumeti	GRUMETI	Lake Victoria	Nyamara	Mara	Serengeti	Natta	MI:
2	2/25/2013	Lottery Club	World vision	Pangani	Majengo	Manyara	Simanjiro	Ngorika	
3	1/28/2013	Unicef	UNICEF	Ruvuma / Southern Coast	Mahakamani	Mtwara	Nanyumbu	Nanyumbu	
4	7/13/2011	Action In A	Artisan	Lake Victoria	Kyanyamisa	Kagera	Karagwe	Nyakasimbi	

5 rows × 28 columns

```
In [48]: df1.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 59400 entries, 0 to 59399
Data columns (total 28 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   date_recorded                        59400 non-null  object
1   funder                              59400 non-null  object
2   installer                           59400 non-null  object
3   basin                              59400 non-null  object
4   subvillage                          59400 non-null  object
5   region                             59400 non-null  object
6   lga                                 59400 non-null  object
7   ward                               59400 non-null  object
8   public_meeting                     59400 non-null  object
9   recorded_by                        59400 non-null  object
10  permit                             59400 non-null  object
11  extraction_type                    59400 non-null  object
12  management                         59400 non-null  object
13  payment                           59400 non-null  object
14  water_quality                     59400 non-null  object
15  quantity                          59400 non-null  object
16  source                            59400 non-null  object
17  waterpoint_type                   59400 non-null  object
18  status_group                      59400 non-null  object
19  amount_tsh                        59400 non-null  float64
20  gps_height                        59400 non-null  float64
21  longitude                         59400 non-null  float64
22  latitude                         59400 non-null  float64
23  num_private                       59400 non-null  float64
24  region_code                       59400 non-null  float64
25  district_code                     59400 non-null  float64
26  population                        59400 non-null  float64
27  construction_year                 59400 non-null  float64
dtypes: float64(9), object(19)
memory usage: 12.7+ MB
```


Checking for duplicate data

```
In [49]: #checking for duplicate data based on all the columns
df[df.duplicated()]
```

[illegible]

	id	amount_tsh	date_recorded	funder	gps_height	installer	longitude	latitude
57662	47039	0.0	10/25/2012	Dwsp	0	DWE	0.0	-2.000000e-08
57807	49622	0.0	8/26/2011	Government Of Tanzania	0	Government	0.0	-2.000000e-08
58463	1562	0.0	2/16/2013	Dwsp	0	DWE	0.0	-2.000000e-08
58859	63207	0.0	10/26/2012	Lwi	0	LWI	0.0	-2.000000e-08
59166	52986	0.0	1/22/2013	World Vision	0	World Vision	0.0	-2.000000e-08

141 rows × 41 columns



We can that there are 141 rows of duplicate data. We can remove them from the dataset

```
In [50]: #removing duplicates
df1.drop_duplicates(inplace=True,keep='first')
df1.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 59259 entries, 0 to 59399
Data columns (total 28 columns):
#   Column                Non-Null Count  Dtype
---  -
0   date_recorded          59259 non-null  object
1   funder                 59259 non-null  object
2   installer              59259 non-null  object
3   basin                 59259 non-null  object
4   subvillage            59259 non-null  object
5   region                59259 non-null  object
6   lga                   59259 non-null  object
7   ward                  59259 non-null  object
8   public_meeting        59259 non-null  object
9   recorded_by           59259 non-null  object
10  permit                 59259 non-null  object
11  extraction_type        59259 non-null  object
12  management             59259 non-null  object
13  payment                59259 non-null  object
14  water_quality          59259 non-null  object
15  quantity               59259 non-null  object
16  source                 59259 non-null  object
17  waterpoint_type        59259 non-null  object
18  status_group           59259 non-null  object
19  amount_tsh             59259 non-null  float64
20  gps_height             59259 non-null  float64
21  longitude              59259 non-null  float64
22  latitude               59259 non-null  float64
23  num_private            59259 non-null  float64
24  region_code            59259 non-null  float64
25  district_code          59259 non-null  float64
26  population             59259 non-null  float64
27  construction_year      59259 non-null  float64
```

dtypes: float64(9), object(19)
memory usage: 13.1+ MB

We now have a dataset that in has no null and duplicate values.

Mapping waterpump distribution


```
In [51]: #create a new df
mapdf = df.copy()

#filter out the Longitude values
mapdf = mapdf[mapdf['longitude'] > 0]

#read the shape file with geopandas
tanzania_map = gpd.read_file('Districts and TC as 2020.shp')
# tanzania_map.plot(color='lightgrey',figsize=(8,8)); just look at the map of tanzania

crs = {'init':'EPSG:4326'} #define CRS
geometry = [Point(xy) for xy in zip(mapdf['longitude'], df['latitude'])] #create Points
geo_df = gpd.GeoDataFrame(mapdf,
                           crs = crs,
                           geometry = geometry) #define the geometry df

#plot the data
fig, ax = plt.subplots(figsize = (10,10))
tanzania_map.to_crs(epsg=4326).plot(ax=ax, color='lightgrey')
sns.scatterplot(x="longitude", y="latitude",data=mapdf,hue='status_group',ax=ax);
ax.set_title('Distribution of Waterpumps in Tanzania');
```



Encoding

```
In [52]: df1_cat = df1.select_dtypes(include=['object'])
df1_cat.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 59259 entries, 0 to 59399
Data columns (total 19 columns):
#   Column                Non-Null Count  Dtype
---  -
0   date_recorded          59259 non-null  object
1   funder                 59259 non-null  object
2   installer              59259 non-null  object
3   basin                 59259 non-null  object
4   subvillage            59259 non-null  object
5   region                59259 non-null  object
6   lga                   59259 non-null  object
7   ward                  59259 non-null  object
8   public_meeting        59259 non-null  object
9   recorded_by           59259 non-null  object
10  permit                 59259 non-null  object
```

```

11 extraction_type 59259 non-null object
12 management     59259 non-null object
13 payment         59259 non-null object
14 water_quality   59259 non-null object
15 quantity        59259 non-null object
16 source          59259 non-null object
17 waterpoint_type 59259 non-null object
18 status_group    59259 non-null object
dtypes: object(19)
memory usage: 9.0+ MB

```

```

In [53]: #Cheking for the number of unique values in each column
cols=[]
unique_number = []
for name in df1_cat.columns:
    cols.append(name)
for item in cols:
    unique_number.append(df1_cat[item].nunique())

unique_dict =dict(zip(cols,unique_number))
unique_dict

```

```

Out[53]: {'date_recorded': 356,
'funder': 1898,
'installer': 2146,
'basin': 9,
'subvillage': 19288,
'region': 21,
'lga': 125,
'ward': 2092,
'public_meeting': 3,
'recorded_by': 1,
'permit': 3,
'extraction_type': 18,
'management': 12,
'payment': 7,
'water_quality': 8,
'quantity': 5,
'source': 10,
'waterpoint_type': 7,
'status_group': 3}

```

Since there are certain features with a large number of unique values, we can try encoding the categorical data by using two different methods to see which works best: label encoding and one hot encoding

Label encoding approach

```

In [54]: #instantiate the encoder
labelencoder = LabelEncoder()

#converting df1_cat into str type
df1_cat = df1_cat.astype('str')

#fit and transform the categrical data
df1_cat_enc = df1_cat.apply(labelencoder.fit_transform)
df1_cat_enc.shape

```

```

Out[54]: (59259, 19)

```

One-hot encoding

```
In [55]: #using the get dummies method to one hot encode
df1_cat_ohe=pd.get_dummies(df1_cat)

#creating another copy of the df for analysis
df3 = df1.copy()
df3.drop(df1_cat.columns,axis=1,inplace=True)

#concatenating the two df's
df3 = pd.concat([df3,df1_cat_ohe],axis=1)
df3.shape
```

Out[55]: (59259, 26021)

We can see that having such a large df would be resource-intensive to work with.

Hence, we will use the LabelEncoding approach for our analysis

```
In [56]: #making a copy for analysis
df2 = df1.copy()

#dropping the categorical columns
df2.drop(df1_cat_enc.columns,axis=1,inplace=True)
df2.info()

#combining the encoded and numeric data
df2 = pd.concat([df2,df1_cat_enc],axis=1)
df2.head()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 59259 entries, 0 to 59399
Data columns (total 9 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   amount_tsh            59259 non-null  float64
 1   gps_height            59259 non-null  float64
 2   longitude             59259 non-null  float64
 3   latitude              59259 non-null  float64
 4   num_private           59259 non-null  float64
 5   region_code           59259 non-null  float64
 6   district_code         59259 non-null  float64
 7   population            59259 non-null  float64
 8   construction_year     59259 non-null  float64
dtypes: float64(9)
memory usage: 4.5 MB
```

```
Out[56]:
```

	amount_tsh	gps_height	longitude	latitude	num_private	region_code	district_code	population	c
0	0.017143	0.517483	0.865979	0.153923	0.0	0.102041	0.0625	0.003574	
1	0.000000	0.520629	0.860047	0.815659	0.0	0.193878	0.0250	0.009180	
2	0.000071	0.271329	0.928504	0.671973	0.0	0.204082	0.0500	0.008197	
3	0.000000	0.123427	0.953922	0.042418	0.0	0.908163	0.7875	0.001902	
4	0.000000	0.031469	0.771612	0.843309	0.0	0.173469	0.0125	0.000000	

5 rows × 28 columns

We now have df2 with encoded values that we can use for modelling.
We will build 3 different baseline models and compare their f1-scores. We will then select the one with the highest f1-score to fine tune and build upon.

Building baseline models

Logistic Regression

```
In [57]: #creating X and y
X=df2.drop('status_group',axis=1)
y=df2[['status_group']]

#splitting the train and test sets
X_train,X_test,y_train,y_test = train_test_split(X,y,random_state=123)

#instantiate logistic regression model
logreg = LogisticRegression(random_state=123)

#fit the model onto the train sets
logreg.fit(X_train,y_train)
logreg

#predict values of the model
y_hat_train = logreg.predict(X_train)
y_hat_test = logreg.predict(X_test)

#evaluate model
names = ['functional','non functional','functional needs repair']
f1_logreg = round(f1_score(y_test,y_hat_test,average='macro'),3)
print(classification_report(y_test,y_hat_test,target_names=names))
```

	precision	recall	f1-score	support
functional	0.54	0.97	0.69	7945
non functional	0.00	0.00	0.00	1091
functional needs repair	0.55	0.06	0.11	5779
accuracy			0.54	14815
macro avg	0.37	0.34	0.27	14815
weighted avg	0.51	0.54	0.42	14815

Decision Tree

```
In [58]: #instantiate
clf = DecisionTreeClassifier(criterion='entropy',random_state=123)

#fit the model onto the train sets
clf.fit(X_train,y_train)

#predict
y_hat_train = clf.predict(X_train)
y_hat_test = clf.predict(X_test)
```

```
#evaluate model
f1_tree = round(f1_score(y_test,y_hat_test,average='macro'),3)
print(classification_report(y_test,y_hat_test,target_names=names))
```

	precision	recall	f1-score	support
functional	0.79	0.79	0.79	7945
non functional	0.38	0.38	0.38	1091
functional needs repair	0.76	0.76	0.76	5779
accuracy			0.75	14815
macro avg	0.64	0.64	0.64	14815
weighted avg	0.75	0.75	0.75	14815

KNN model

```
In [59]: #instantiate
knn_baseline_model = KNeighborsClassifier()

#fit onto the data
knn_baseline_model.fit(X_train,y_train)

#predict
y_hat_train = knn_baseline_model.predict(X_train)
y_hat_test = knn_baseline_model.predict(X_test)

#evaluate model
f1_knn = round(f1_score(y_test,y_hat_test,average='macro'),3)
print(classification_report(y_test,y_hat_test,target_names=names))
```

	precision	recall	f1-score	support
functional	0.64	0.77	0.70	7945
non functional	0.30	0.15	0.20	1091
functional needs repair	0.59	0.49	0.54	5779
accuracy			0.61	14815
macro avg	0.51	0.47	0.48	14815
weighted avg	0.60	0.61	0.60	14815

Selecting a model

```
In [60]: #import the metrics library
print(f'f1-score of baseline logistic regression is {f1_logreg}')
print(f'f1-score of baseline decision tree is {f1_tree}')
print(f'f1-score of knn model is {f1_knn}')
```

```
f1-score of baseline logistic regression is 0.269
f1-score of baseline decision tree is 0.644
f1-score of knn model is 0.477
```

Since, the decision tree baseline model has the highest performance score, we will build on that for further analysis

Decision Tree

```
In [61]: #re-run the baseline model so all variables are updated correctly
```

```

#for the train and test sets

#instantiate
clf = DecisionTreeClassifier(criterion='entropy',random_state=123)

#fit the model onto the train sets
clf.fit(X_train,y_train)
clf

#predict
y_hat_train = clf.predict(X_train)
y_hat_test = clf.predict(X_test)

#evaluate model
from sklearn.metrics import classification_report

print('TRAIN SCORES')
print(classification_report(y_train,y_hat_train,target_names=names))
print('-----')
print('TEST SCORES')
print(classification_report(y_test,y_hat_test,target_names=names))

accuracy_baseline_train = accuracy_score(y_train,y_hat_train)
accuracy_baseline_test = accuracy_score(y_test,y_hat_test)

```

```

TRAIN SCORES

```

	precision	recall	f1-score	support
functional	1.00	1.00	1.00	24241
non functional	1.00	1.00	1.00	3217
functional needs repair	1.00	1.00	1.00	16986
accuracy			1.00	44444
macro avg	1.00	1.00	1.00	44444
weighted avg	1.00	1.00	1.00	44444

```

-----
TEST SCORES

```

	precision	recall	f1-score	support
functional	0.79	0.79	0.79	7945
non functional	0.38	0.38	0.38	1091
functional needs repair	0.76	0.76	0.76	5779
accuracy			0.75	14815
macro avg	0.64	0.64	0.64	14815
weighted avg	0.75	0.75	0.75	14815

We can see that the model is clearly overfitting on the train dataset.

Hyperparameter Tuning

max_depth

```

In [62]: #creating a list of depth values
max_depth = np.arange(1,50)

#creating an empty list to store scores for each depth
f1_test_scores = []

```



```

# create a loop for the classifier to run with the different depth values
for depth in max_depth:

    #instantiate
    classifier = DecisionTreeClassifier(criterion='entropy',max_depth=depth,random_stat

    #fit the model
    classifier.fit(X_train,y_train)

    #predict values
    y_hat_test = classifier.predict(X_test)


    #calculate
    f1 = f1_score(y_test,y_hat_test,average='macro')

    #add the scores to the list
    f1_test_scores.append(f1)

#visualize the data
import matplotlib.pyplot as plt
fig,ax = plt.subplots(figsize=(7,7))
ax.plot(max_depth,f1_test_scores,c='b')
ax.set_xlabel('max_depth')
ax.set_ylabel('f1-score')
ax.set_title('max_depth')

```

Out[62]: Text(0.5, 1.0, 'max_depth')



We can see that the accuracy for the model peaks at 20 before decreasing and

levelling off

min_samples_split

```
In [63]: #use the optimum value of depth
depth = 20

#define a range of min_samples_for each split
min_samples_range = np.arange(18,30)

#create a loop with the optimum depth and different min_samples
#creating an empty list to store scores for each depth
f1_test_scores = []
# accuracy_scores = []
# precision_scores = []
# recall_scores = []

# create a loop for the classifier to run with the different depth values
for sample in min_samples_range:

    #instantiate
    classifier = DecisionTreeClassifier(criterion='entropy',max_depth=depth,min_samples

    #fit the model
    classifier.fit(X_train,y_train)


    #predict values
    y_hat_test = classifier.predict(X_test)

    #calculate
    f1 = f1_score(y_test,y_hat_test,average='macro')

    #add the scores to the list
    f1_test_scores.append(f1)

#visualize
fig,ax = plt.subplots(figsize=(7,7))
ax.plot(min_samples_range,f1_test_scores,c='b')

ax.set_xlabel('min_samples')
ax.set_ylabel('f1-score')
ax.set_title('min_samples_split')
plt.show();
```



We can see that the accuracy peaks at a sample value of 21

min_samples_leaf

```
In [64]: #use the optimum value of depth,min_samples_split
depth = 20
min_samples = 21

#define a range of min_samples_for each split
min_leaf_range = np.arange(2,20)

#create a loop with the optimum depth and different min_samples
#creating an empty list to store scores for each depth
f1_test_scores = []

# create a loop for the classifier to run with the different depth values
for sample in min_leaf_range:

    #instantiate
    classifier = DecisionTreeClassifier(criterion='entropy',max_depth=depth,min_samples
                                       min_samples_leaf=sample, random_state=123)

    #fit the model
    classifier.fit(X_train,y_train)

    #predict values
    y_hat_test = classifier.predict(X_test)

    #calculate
    f1 = f1_score(y_test,y_hat_test,average='macro')
```

```


#add the scores to the list
f1_test_scores.append(f1)

#visual
fig,ax = plt.subplots(figsize=(7,7))
ax.plot(min_leaf_range,f1_test_scores,c='b')

ax.set_xlabel('min_leaf')
ax.set_ylabel('f1-score')
ax.set_title('min_leaf_samples')

plt.show();

```



We can see that the peak value is 3

model with optimized parameters

```

In [65]: #run the model with the optimized parameters
depth=20
min_samples=21
leaf=3

#instantiate
classifier = DecisionTreeClassifier(criterion='entropy',max_depth=depth,min_samples_split=min_samples,
min_samples_leaf=leaf,random_state=123)

#fit the model
classifier.fit(X_train,y_train)

#predict values

```

```

y_hat_train = classifier.predict(X_train)
y_hat_test = classifier.predict(X_test)

#scores
print('TRAIN SCORES')
print('-----')
print(classification_report(y_train,y_hat_train,target_names=names))
print('TEST SCORES')
print('-----')
print(classification_report(y_test,y_hat_test,target_names=names))

accuracy_optimized_train = accuracy_score(y_train,y_hat_train)
accuracy_optimized_test = accuracy_score(y_test,y_hat_test)

f1_score_optimized_train = f1_score(y_train,y_hat_train,average='macro')
f1_score_optimized_test = f1_score(y_test,y_hat_test,average='macro')

```

TRAIN SCORES

```

-----
              precision    recall  f1-score   support

functional      0.85        0.92        0.89       24241
non functional   0.71        0.49        0.58        3217
functional needs repair  0.89        0.83        0.86       16986

accuracy              0.86       44444
macro avg      0.82        0.75        0.77       44444
weighted avg   0.86        0.86        0.85       44444

```

TEST SCORES

```

-----
              precision    recall  f1-score   support

functional      0.78        0.85        0.81        7945
non functional   0.47        0.34        0.39       1091
functional needs repair  0.79        0.74        0.76       5779

accuracy              0.77       14815
macro avg      0.68        0.64        0.66       14815
weighted avg   0.76        0.77        0.76       14815

```

feature_importance

```

In [66]: #creating a df with just feature_importance data
df_importance = pd.DataFrame({'features':X_train.columns,
                              'importance':classifier.feature_importances_})

df_importance.head()

```

```

Out[66]:
   features  importance
0  amount_tsh    0.027697
1   gps_height    0.042343
2   longitude    0.113452
3    latitude    0.090703
4  num_private    0.000268


```

```
In [67]: #sorting the importance in ascending order
df_importance_sorted = df_importance.sort_values(by=['importance'],ascending=True)
df_importance_sorted.head()
```

```
Out[67]:
```

	features	importance
18	recorded_by	0.000000
4	num_private	0.000268
23	water_quality	0.004597
19	permit	0.004823
17	public_meeting	0.005430

```
In [68]: #plot the data
fig,ax=plt.subplots(figsize=(10,10))
df_importance_sorted.plot(kind='barh',ax=ax);
ax.set_yticklabels(df_importance_sorted['features'])
ax.set_title('feature_importance');
ax.legend(loc=4);
```



Let's build a model based on the top_10 features to see if we can get better

results

extracting top_10 features

```
In [69]: #create a new df with the top_10 features alone

#get top10 features as a df
cols = df_importance_sorted['features'].tail(10)
cols=pd.DataFrame(data=cols)
cols
```

```
Out[69]:
```

	features
15	lga
1	gps_height
13	subvillage
8	construction_year
9	date_recorded
10	funder
26	waterpoint_type
3	latitude
2	longitude
24	quantity

```
In [70]: #df with top10 features
df_top10 =df2[list(cols['features'])]

#combine with the target variable
df_top10 = pd.concat([df_top10,df1['status_group']],axis=1)

print(df_top10.shape)
df_top10.head()
```

(59259, 11)

```
Out[70]:
```

	lga	gps_height	subvillage	construction_year	date_recorded	funder	waterpoint_type	latitude	lon
0	51	0.517483	11808	0.993045	171	1370	1	0.153923	0
1	103	0.520629	15839	0.998510	216	469	1	0.815659	0
2	108	0.271329	9075	0.998013	144	825	2	0.671973	0
3	87	0.123427	8983	0.986587	21	1742	2	0.042418	0
4	26	0.031469	7698	0.000000	268	20	1	0.843309	0

Decision Tree with the top10 features

GridSearch CV

Rather than use the earlier approach, where we built a baseline model and then tuned each hyper parameter separately, we can combine all these steps into one using GridSearchCV

```
In [71]: #define X & y
X = df_top10.drop('status_group',axis=1)
y=df_top10[['status_group']]

#split the data
X_top10_train,X_top10_test,y_top10_train,y_top10_test = train_test_split(X,y,random_state=123)

#instantiate
clf_top10 = DecisionTreeClassifier(random_state=123)

#define the parameter grid
param_grid = {'max_depth': np.arange(20,25),
              'min_samples_split': np.arange(20,25),
              'min_samples_leaf': np.arange(3,7)
              }

#instantiate
gs_tree = GridSearchCV(estimator=clf_top10,param_grid=param_grid,cv=5)

#fit
gs_tree.fit(X_top10_train,y_top10_train)

#predict
gs_tree_train = gs_tree.predict(X_top10_train)
gs_tree_test = gs_tree.predict(X_top10_test)

print('TRAIN SCORES')
print('-----')
print(classification_report(y_top10_train,gs_tree_train))
print('TEST SCORES')
print('-----')
print(classification_report(y_top10_test,gs_tree_test))

accuracy_grid_train = accuracy_score(y_top10_train,gs_tree_train)
accuracy_grid_test = accuracy_score(y_top10_test,gs_tree_test)

f1_score_gs_tree_train = f1_score(y_top10_train,gs_tree_train,average='macro')
f1_score_gs_tree_test = f1_score(y_top10_test,gs_tree_test,average='macro')
```

TRAIN SCORES

	precision	recall	f1-score	support
functional	0.83	0.92	0.87	24241
functional needs repair	0.67	0.40	0.50	3217
non functional	0.86	0.80	0.83	16986
accuracy			0.84	44444
macro avg	0.79	0.71	0.73	44444
weighted avg	0.83	0.84	0.83	44444

TEST SCORES

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

functional	0.77	0.86	0.81	7945
functional needs repair	0.49	0.30	0.37	1091
non functional	0.79	0.72	0.75	5779
accuracy			0.76	14815
macro avg	0.68	0.63	0.65	14815
weighted avg	0.76	0.76	0.76	14815

```
In [72]: #getting the best parameters
gs_tree.best_params_
```

```
Out[72]: {'max_depth': 20, 'min_samples_leaf': 5, 'min_samples_split': 23}
```

Random Forest with top10 features

Let's use a popular ensemble method called Random Forest on the top_10 features to and compare it to the GridSearchCV model to see if we can make improvements. Random Forest combines Bootstrapping and Sub-Space Sampling methods to build models that are robust and immune to noise in the data.

```
In [73]: #instantiante the classifier with the same parameters from earlier
forest=RandomForestClassifier(n_estimators=100,max_depth=depth,min_samples_split=min_sa
                                min_samples_leaf=leaf)

#fit the data
forest.fit(X_top10_train,y_top10_train)

#predict
forest_y_train = forest.predict(X_top10_train)
forest_y_test = forest.predict(X_top10_test)

#evaluate

# print('TRAIN SCORES')
# print('-----')
# print(classification_report(y_top10_train,forest_y_train))
print('TEST SCORES')
print('-----')
print(classification_report(y_top10_test,forest_y_test))

accuracy_forest_train = accuracy_score(y_top10_train,forest_y_train)
accuracy_forest_test = accuracy_score(y_top10_test,forest_y_test)

f1_score_rf_train = f1_score(y_top10_train,forest_y_train,average='macro')
f1_score_rf_test = f1_score(y_top10_test,forest_y_test,average='macro')
```

TEST SCORES


	precision	recall	f1-score	support
functional	0.77	0.91	0.83	7945
functional needs repair	0.68	0.23	0.34	1091
non functional	0.84	0.73	0.78	5779
accuracy			0.79	14815
macro avg	0.76	0.62	0.65	14815

weighted avg 0.79 0.79 0.78 14815

Visualize scores

```
In [74]: # create a dataframe of the scores for the different models
score_df = pd.DataFrame({'train_accuracy': [accuracy_optimized_train, accuracy_grid_train, accuracy_forest_train],
                        'test_accuracy': [accuracy_optimized_test, accuracy_grid_test, accuracy_forest_test],
                        'train_f1': [f1_score_optimized_train, f1_score_grid_train, f1_score_forest_train],
                        'test_f1': [f1_score_optimized_test, f1_score_grid_test, f1_score_forest_test]},
                        index=['optimized', 'grid', 'forest'])

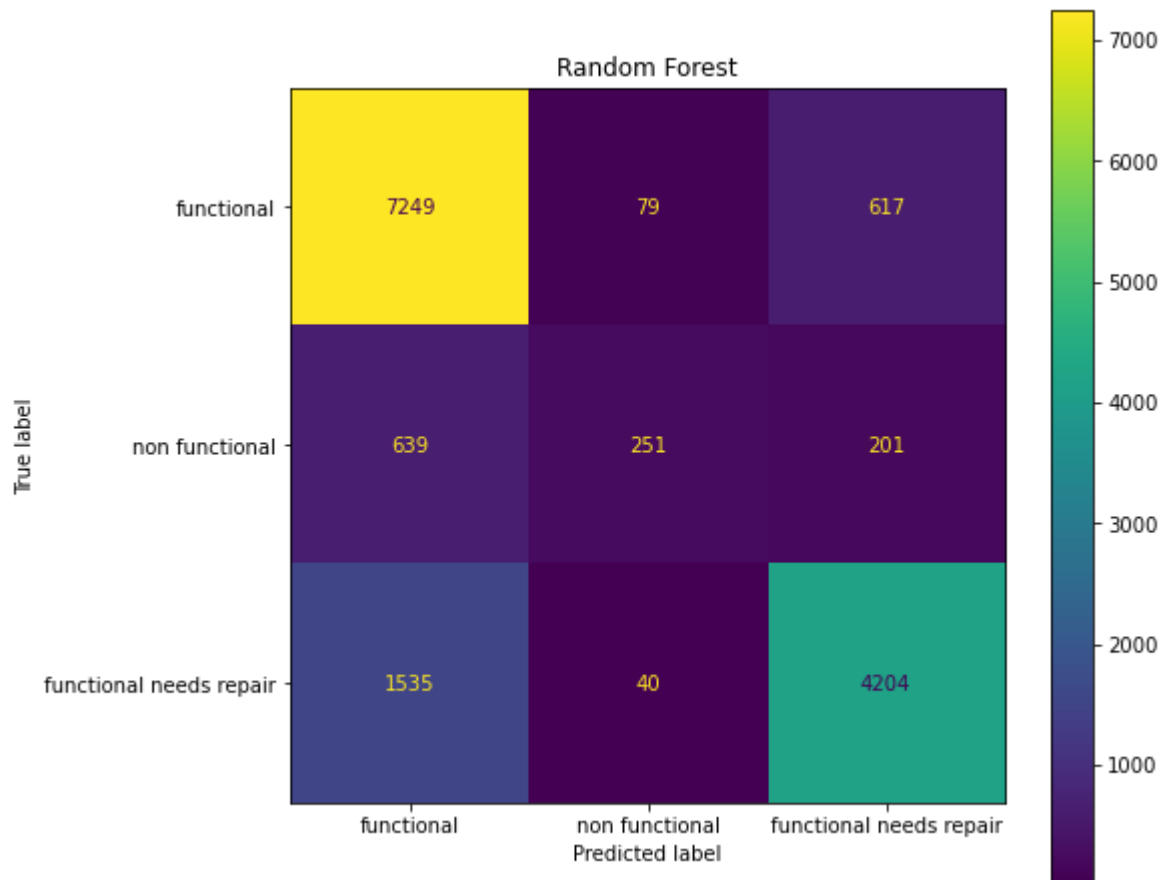
#plot the data
fig, ax = plt.subplots(figsize=(8, 8))
score_df.plot(kind='line', ax=ax, marker='x');
```



Confusion Matrix

```
In [75]: #Plot the confusion matrix of the random forest model
fig, ax = plt.subplots(figsize=(8, 8))
plot_confusion_matrix(forest, X_top10_test, y_top10_test, ax=ax, display_labels=names);
ax.set_title('Random Forest')
```

Out[75]: Text(0.5, 1.0, 'Random Forest')



We can clearly see that the 'functional needs repair' class is a problem for the model

Earlier, we built the model with only the top10 features based on feature importance. Let's now try and build a model with all the features and see if more data helps improve performance.

Random Forest with all the features

```
In [76]: #define X,y
X=df2.drop('status_group',axis=1)
y=df2[['status_group']]

#using the cleaned dataset with all features
X_train,X_test,y_train,y_test = train_test_split(X,y,random_state=123)

#instantiate the classifier
new_forest =RandomForestClassifier(n_estimators=100,max_depth=depth,min_samples_split=m
min_samples_leaf=leaf)

#fit the data
new_forest.fit(X_train,y_train)

#predict
y_hat_train = new_forest.predict(X_train)
y_hat_test = new_forest.predict(X_test)

#scores
# print('TRAIN SCORES')
```


```
# print('-----')
# print(classification_report(y_train,y_hat_train,target_names=names))
print('TEST SCORES')
print('-----')
print(classification_report(y_test,y_hat_test,target_names=names))
```

TEST SCORES

	precision	recall	f1-score	support
functional	0.78	0.92	0.84	7945
non functional	0.72	0.25	0.37	1091
functional needs repair	0.86	0.75	0.80	5779
accuracy			0.80	14815
macro avg	0.78	0.64	0.67	14815
weighted avg	0.81	0.80	0.79	14815

Confusion matrix

```
In [77]: fig,(ax1,ax2)=plt.subplots(figsize=(20,10),nrows=1,ncols=2)
plot_confusion_matrix(forest,X_top10_test,y_top10_test,ax=ax1,display_labels=names)
plot_confusion_matrix(new_forest,X_test,y_test,ax=ax2,display_labels=names);
fig.tight_layout(pad=10.0)
ax1.set_title('top10 features only');
ax2.set_title('all features');
```



We can see that it makes a slight difference to the model.

Examining the target feature

```
In [78]: #examining the target variable
df1['status_group'].value_counts()
```

```
Out[78]: functional      32186
non functional    22765
functional needs repair    4308
Name: status_group, dtype: int64
```

Clearly, we can see an imbalance between the classes. We can try to train a

model with equal representation from each class and check the results.

Training a model with a balanced data set

```
In [79]: #seperate each class into a seperate class with the same number of rows as the repair c
functional = df2[df2['status_group'] == 0]
functional = functional.iloc[0:4308,:]

non_functional = df2[df2['status_group'] == 2]
non_functional = non_functional.iloc[0:4308,:]

repair = df2[df2['status_group'] == 1]
repair['status_group'].value_counts()

#concatenate all three df's
new_df = pd.concat([functional,non_functional,repair])

#check the value counts
new_df['status_group'].value_counts()
```

```
Out[79]: 2    4308
1    4308
0    4308
Name: status_group, dtype: int64
```

```
In [80]: #building the baseline model

#define X,y
X=new_df.drop('status_group',axis=1)
y=new_df[['status_group']]

#using the cleaned dataset with all features
X_train_new,X_test_new,y_train_new,y_test_new = train_test_split(X,y,random_state=123)

#instantiate the classifier
newdf_forest =RandomForestClassifier()

#fit the data on the new_df
forest1=newdf_forest.fit(X_train_new,y_train_new)

#predict score on the new_df
y_hat_train_new = new_forest.predict(X_train_new)
y_hat_test_new = new_forest.predict(X_test_new)

#scores
# print('TRAIN SCORES')
# print('-----')
# print(classification_report(y_train_new,y_hat_train_new,target_names=names))
print('TEST SCORES')
print('-----')
print(classification_report(y_test_new,y_hat_test_new,target_names=names))
```

TEST SCORES

	precision	recall	f1-score	support
functional	0.56	0.95	0.70	1074
non functional	0.97	0.29	0.44	1067
functional needs repair	0.82	0.81	0.82	1090

accuracy			0.68	3231
macro avg	0.78	0.68	0.65	3231
weighted avg	0.78	0.68	0.65	3231

```
In [81]: # with the model trained on the balanced dataset, let's check for predictions on the un
y_preds_train = forest1.predict(X_train) # original df
y_preds_test = forest1.predict(X_test) #original df

print('TEST SCORES')
print('-----')
print(classification_report(y_test,y_preds_test,target_names=names))
```

TEST SCORES


```
-----
              precision    recall  f1-score   support

   functional         0.87      0.70      0.77       7945
  non functional         0.34      0.95      0.50       1091
functional needs repair         0.82      0.76      0.79       5779

   accuracy         0.74
  macro avg         0.68
 weighted avg         0.81
```

Confusion Matrix

```
In [82]: fig,(ax1,ax2)=plt.subplots(figsize=(20,10),nrows=1,ncols=2)
plot_confusion_matrix(new_forest,X_test,y_test,ax=ax1,display_labels=names)
plot_confusion_matrix(forest1,X_test,y_test,ax=ax2,display_labels=names);
fig.tight_layout(pad=10.0)
ax1.set_title('trained on unbalanced data');
ax2.set_title('trained on balanced data');
```




Random forest model on the balanced dataset

```
In [83]: #instantaiate with optimized parameters from before
forest2 = RandomForestClassifier(max_depth=depth,min_samples_split=min_samples,
                                min_samples_leaf=leaf,random_state=42)
```

```
#fit the data on the balanced dataset
forest2.fit(X_train_new,y_train_new)

# predict data on the original df
y_hat_test_preds = forest2.predict(X_test) # original df
```

```
In [84]: fig,ax=plt.subplots(figsize=(8,8))
plot_confusion_matrix(forest2,X_test,y_test,display_labels=names,ax=ax);
ax.set_title('Random Forest on balanced data');
```



```
In [85]: print('TEST SCORES')
print('-----')
print(classification_report(y_test,y_hat_test_preds,target_names=names))
```

TEST SCORES

```
-----
              precision    recall  f1-score   support

   functional         0.84      0.68      0.75       7945
  non functional         0.29      0.83      0.43       1091
functional needs repair         0.81      0.72      0.76       5779

   accuracy              0.71      14815
  macro avg              0.64      0.75      0.65      14815
 weighted avg              0.79      0.71      0.73      14815
```

Next Steps

1. Possibly re-frame this as a binary classification problem i.e functional vs non-functional and see if we can build a better model.
2. Re-create the model with equal number of data points between functional and non-functional. Optimize parameters on this balanced dataset and test it on validation data to check for performance.